



Dr. Hakim Mellah

**Department of Computer Science & Software Engineering
Concordia University, Montreal, Canada**

Lecture 12: I/O System

COMP 346: Operating Systems

These slides has been extracted, modified and updated from original slides of :

- *Operating System Concepts, 10th Edition*, by: Silberschatz/Galvin/Gagne, published by John Wiley & Sons

Overview

- I/O management is a major component of operating system design and operation
 - ❖ Important aspect of computer operation
 - ❖ I/O devices vary greatly
 - ❖ Various methods to control them
 - ❖ Performance management
 - ❖ New types of devices frequent
- Ports, busses, device controllers connect to various devices
- Device drivers encapsulate device details
 - ❖ Present uniform device-access interface to I/O subsystem

I/O Hardware


➤ Incredible variety of I/O devices

- ❖ Storage
- ❖ Transmission
- ❖ Human-interface

➤ Common concepts – signals from I/O devices interface with computer

- ❖ **Port** – connection point for device
- ❖ **Bus** - **daisy chain** or shared direct access
 - ✓ **PCI** bus common in PCs and servers, PCI Express (**PCIe**)
 - ✓ **expansion bus** connects relatively slow devices
- ❖ **Controller** (**host adapter**) – electronics that operate port, bus, device
 - ✓ Sometimes integrated
 - ✓ Sometimes separate circuit board (host adapter)
 - ✓ Contains processor, microcode, private memory, bus controller, etc

A Typical PC Bus Structure



I/O Hardware (Cont.)

- I/O instructions control devices
- Devices usually have registers where device driver places commands, addresses, and data to write, or read data from registers after command execution
 - ❖ Data-in register, data-out register, status register, control register
 - ❖ Typically 1-4 bytes, or FIFO buffer
- Devices have addresses, used by
 - ❖ Direct I/O instructions
 - ❖ **Memory-mapped I/O**
 - ✓ Device data and command registers mapped to processor address space
 - ✓ Especially for large address spaces (graphics)

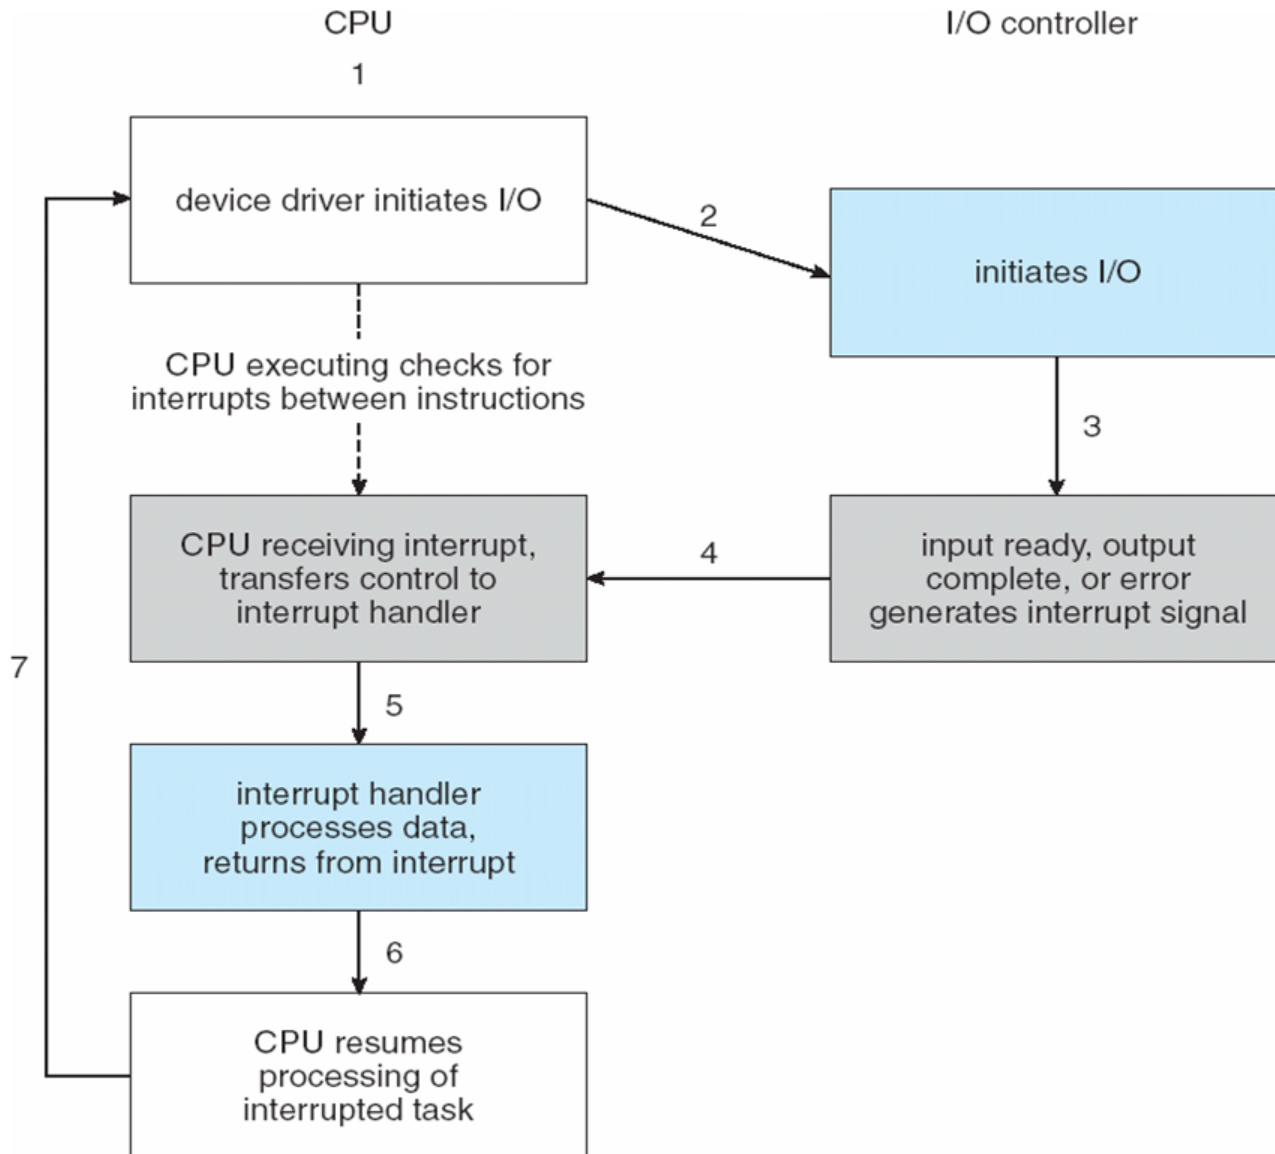
Polling

- ❑ For each byte of I/O
 1. Read busy bit from status register until 0
 2. Host sets read or write bit and if write copies data into data-out register
 3. Host sets command-ready bit
 4. Controller sets busy bit, executes transfer
 5. Controller clears busy bit, error bit, command-ready bit when transfer done
- ❑ Step 1 is **busy-wait** cycle to wait for I/O from device
 - ❑ Reasonable if device is fast
 - ❑ But inefficient if device slow
 - ❑ CPU switches to other tasks?
 - ▶ But if miss a cycle data overwritten / lost

Interrupts

- Polling can happen in 3 instruction cycles
 - ❖ Read status, logical-and to extract status bit, branch if not zero
- CPU **Interrupt-request line** triggered by I/O device
 - ❖ Checked by processor after each instruction
- **Interrupt handler** receives interrupts
 - ❖ **Maskable** to ignore or delay some interrupts
- **Interrupt vector** to dispatch interrupt to correct handler
 - ❖ Context switch at start and end
 - ❖ Based on priority
 - ❖ Some **nonmaskable**
 - ❖ Interrupt chaining if more than one device at same interrupt number

Interrupt-Driven I/O Cycle



Intel Pentium Processor Event-Vector Table

vector number	description
0	divide error
1	debug exception
2	null interrupt
3	breakpoint
4	INTO-detected overflow
5	bound range exception
6	invalid opcode
7	device not available
8	double fault
9	coprocessor segment overrun (reserved)
10	invalid task state segment
11	segment not present
12	stack fault
13	general protection
14	page fault
15	(Intel reserved, do not use)
16	floating-point error
17	alignment check
18	machine check
19–31	(Intel reserved, do not use)
32–255	maskable interrupts

Interrupts (Cont.)

- Interrupt mechanism also used for **exceptions**
 - ❖ Terminate process, crash system due to hardware error
- Page fault executes when memory access error
- System call executes via **trap** to trigger kernel to execute request
- Multi-CPU systems can process interrupts concurrently
 - ❖ If operating system designed to handle it
- Used for time-sensitive processing, frequent, must be fast

Direct Memory Access

- Used to avoid **programmed I/O** (one byte at a time) for large data movement
- Requires **DMA** controller
- Bypasses CPU to transfer data directly between I/O device and memory
- OS writes DMA command block into memory
 - ❖ Source and destination addresses
 - ❖ Read or write mode
 - ❖ Count of bytes
 - ❖ Writes location of command block to DMA controller
 - ❖ Bus mastering of DMA controller – grabs bus from CPU
 - ✓ **Cycle stealing** from CPU but still much more efficient
 - ❖ When done, interrupts to signal completion
- Version that is aware of virtual addresses can be even more efficient - **DVMA**

Six Step Process to Perform DMA Transfer

